

Elements Of Data Science - S2022

Week 7: Model Evaluation and Hyperparameter Tuning

3/1/2022

TODOs

- Readings:
 - **PDSH Chapter 5: In-Depth: Decision Trees and Random Forests**
 - HOML: Chapter 6
 - HOML: Chapter 7
 - **PDSH Chapter 5: Feature Engineering**
- **Quiz 7**, due **Monday March 7th, 11:59pm ET**
- **HW2**, out this week **Fri March 11th, 11:59pm ET**
- **Project**
 - March 22nd presentation
 - Details will be released on Canvas

Today

- Model Evaluation and Selection
- Hyperparameter Tuning
- Regularization

Questions?

Environment Setup

In [1]:

```
import numpy
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from mlxtend.plotting import plot_decision_regions

sns.set_style('darkgrid')

%matplotlib inline
```

Model Evaluation and Hyperparameter Tuning

- How well are any of our models working?
- How can we compare different models?
- How do we decide on hyperparameter settings?
- How can we keep our models from "overfitting"?
- How do we do all this both for Regression and Classification?

How well are our models performing?

Regression

- Mean Squared Error (MSE) and Root Mean Squared Error (RMSE)
- R^2 or R^2
- (Adjusted R^2 - will talk about during Feature Selection)

Classification

- Accuracy
- Precision/Recall/F1
- ROC Area Under the Curve (AUC)

Data Setup for Regression

In [2]:

```
zscore = lambda x: (x - x.mean()) / x.std()

df_wine = pd.read_csv('../data/wine_dataset.csv',
                      usecols=['alcalinity_of_ash', 'magnesium', 'alcohol', 'ash', 'proline', 'hue', 'class'])
numeric_cols = ['alcalinity_of_ash', 'magnesium', 'alcohol', 'ash', 'proline', 'hue']

df_wine[numeric_cols] = df_wine[numeric_cols].apply(zscore) # standardize numeric feature cols

X = df_wine[['proline', 'hue', 'ash']]

y_r = df_wine['alcohol'] # regression target
```

In [3]:

```
df_wine.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 178 entries, 0 to 177
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	alcohol	178 non-null	float64
1	ash	178 non-null	float64
2	alcalinity_of_ash	178 non-null	float64
3	magnesium	178 non-null	float64

4	hue	178	non-null	float64
5	proline	178	non-null	float64
6	class	178	non-null	int64

dtypes: float64(6), int64(1)
memory usage: 9.9 KB


Regression with Simple Linear Model

In [4]:

```
from sklearn.linear_model import LinearRegression

lr = LinearRegression().fit(X[['proline']],y_r)
argmax_y_r = np.argmax(y_r)
y_pred = lr.predict(X[['proline']])

fig,ax = plt.subplots(1,1,figsize=(8,6))
sns.scatterplot(x=X.proline, y=y_r, color='grey');
ax.plot(X.proline,y_pred,color='k');
ax.vlines(X.proline.iloc[argmax_y_r],y_r.iloc[argmax_y_r],y_pred[argmax_y_r],color='r');
```



How Good is This Fit? MSE and RMSE

- **Mean Squared Error:** $\frac{1}{n} \sum_i (y_i - \hat{y}_i)^2$

In [5]:

```
from sklearn.metrics import mean_squared_error

lr_mse = mean_squared_error(y_r, y_pred)
print(f'{lr_mse = :0.2f}')
```

lr_mse = 0.58

- But this is the squared error! (alcohol²)
- **Root Mean Squared Error:** $\sqrt{\frac{1}{n} \sum_i (y_i - \hat{y}_i)^2}$

In [6]:

```
def root_mean_squared_error(y_true, y_pred):
    return np.sqrt(mean_squared_error(y_true, y_pred))

lr_rmse = root_mean_squared_error(y_r, y_pred)
print(f'{lr_rmse = :0.2f}')
```

lr_rmse = 0.76

Is this good? Need a Baseline Comparison

- What's a baseline to compare against?
- Simple one for Regression: always predict the mean of the targets

In [7]:

```
from sklearn.dummy import DummyRegressor

dummyr = DummyRegressor(strategy='mean') # default strategy
dummyr.fit(X[['proline']], y_r)

dummy_rmse = root_mean_squared_error(y_r, dummyr.predict(X[['proline']]))

print(f'{dummy_rmse = :0.2f}')
```

`dummy_rmse = 1.00`

Comparing against the mean: R^2

- the proportion of variance explained by the model

$$R^2 = 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2}$$

- maximum value of 1
- a value below 0 means the model is predicting worse than just predicting the mean
- sklearn uses R^2 as the default for regression scoring

In [8]:

```
r2_lr = lr.score(X[['proline']], y_r)
r2_dummyr = dummyr.score(X[['proline']], y_r)

print(f'r2_dummyr = :0.2f\nr2_lr = :0.2f')
```

r2_dummyr = 0.00

r2_lr = 0.41

Can we do better?

In [9]:

```
from sklearn.tree import DecisionTreeRegressor


dtr = DecisionTreeRegressor(max_depth=10)
dtr.fit(X[['proline']],y_r)
r2_dtr = dtr.score(X[['proline']],y_r)
print(f'{r2_lr = :0.2f}\n{r2_dtr = :0.2f}')
```

r2_lr = 0.41

r2_dtr = 0.76

In [10]:

```
X_proline_sorted = X[['proline']].sort_values(by='proline').values
X_query = np.linspace(X_proline_sorted.min(),X_proline_sorted.max(),1000).reshape(-1,1)
y_pred = dtr.predict(X_query)
fig,ax = plt.subplots(1,1,figsize=(8,6))
sns.scatterplot(x=X.proline, y=y_r,color='gray');
ax.plot(X_query,y_pred,color='b');
```



But is this what we want? Interpretation vs Prediction

Always good to ask:

- do we want our model to very closely fit our data for interpretation?
- do we want our model to predict well on new, unseen data?

Generalization:


- how well will model predict on data that it hasn't seen yet?

But we used all of our data to train?

- Need to do a **Train/Test Split** to create a held-aside set

Train/Test Split

- **Training Set:** portion of dataset used for training
- **Test/Held-Aside/Out of sample:** portion of dataset used for evaluation
- Want the test set to reflect the same distribution as training



From https://www.researchgate.net/figure/Train-Test-Data-Split_fig6_325870973

Train/Test split with Sklearn

In [11]:

```
from sklearn.model_selection import train_test_split

X_train_r, X_test_r, y_train_r, y_test_r = train_test_split(X[['proline']],
                                                            y_r,
                                                            train_size=.75, #default (only need one of train/test)
                                                            test_size=.25, #default
                                                            random_state=123)

print(f'n_total = {X[['proline']].shape[0]}')
print(f'n_train = {X_train_r.shape[0]}')
print(f'n_test = {X_test_r.shape[0]}')
```

```
n_total = 178
n_train = 133
n_test  = 45
```

- How big should test be?
 - Large enough to capture variance of dataset.
 - Depends on the dataset and the models being trained

Training and Evaluate on Different Data

In [12]:

```
dummyr = DummyRegressor().fit(X_train_r,y_train_r)
lr = LinearRegression().fit(X_train_r,y_train_r)
dtr = DecisionTreeRegressor(max_depth=10).fit(X_train_r,y_train_r)

r2_dummyr = dummyr.score(X_test_r,y_test_r)
r2_lr = lr.score(X_test_r,y_test_r)
r2_dtr = dtr.score(X_test_r,y_test_r)

print(f'{r2_dummyr = : 0.2f}\n{r2_lr = : 0.2f}\n{r2_dtr = : 0.2f}')
```

r2_dummyr = -0.03

r2_lr = 0.28


r2_dtr = -0.31

- DecisionTree model is doing worse than the Dummy model on the test set!

Overfitting and Underfitting


- **Overfitting:** poor generalization due to complexity
 - learning noise in training data
- **Underfitting:** poor generalization due to simplicity
 - not flexible enough to learn concept
- Need to find a balance between simplicity and complexity
- Need to find a balance between **bias** and **variance**

Bias-Variance Tradeoff



From PDSH

Bias-Variance Tradeoff



From PDSH

Bias-Variance Tradeoff Continued

- How close is the model to the underlying concept?
- How sensitive is the model to the training set?





Fig. 1: Graphical Illustration of bias-variance trade-off , Source: Scott Fortmann-Roe., Understanding Bias-Variance Trade-off

Bias-Variance Tradeoff Continued



- We'd like to:
 - reduce the Bias (use a model complex enough to capture the concept)
 - without introducing too much Variance (overfit the data)
 - all in order to minimize **Generalization Error**

Overfitting/Underfitting Revisited


- **Overfitting:** poor generalization due to complexity
 - learning noise in training data
 - model has **high variance and low bias**
- **Underfitting:** poor generalization due to simplicity
 - not flexible enough to learn concept
 - model has **high bias and low variance**

Avoiding Overfitting/Underfitting

- Never train and evaluate on the same set of data!
 - train test split
 - **cross-validation**
- Keep the model as simple as possible (Occom's Razor)

In [13]:


```
fig,ax = plt.subplots(1,3,figsize=(16,4))
for i in range(3):
    sns.scatterplot(x=X_train_r.proline,y=y_train_r,color="gray",ax=ax[i])
    sns.scatterplot(x=X_test_r.proline,y=y_test_r,color="black",s=100,ax=ax[i]);
ax[0].plot(X_query,dummyr.predict(X_query),color='r');
ax[1].plot(X_query,lr.predict(X_query),color='g');
ax[2].plot(X_query,dtr.predict(X_query),color='b');
```



Overfitting? Simplify the model

In [14]:

```
max_depths = [2,3,4]
fig,ax = plt.subplots(1,3,figsize=(16,4))
for i in range(3):
    dtr = DecisionTreeRegressor(max_depth=max_depths[i]).fit(X_train_r,y_train_r)
    sns.scatterplot(x=X_train_r.proline,y=y_train_r,color="gray",ax=ax[i])
    sns.scatterplot(x=X_test_r.proline,y=y_test_r,color="black",s=100,ax=ax[i]);
    ax[i].plot(X_query,dtr.predict(X_query),color='b');
    ax[i].set_title(f'depth {max_depths[i]}, r2: {dtr.score(X_test_r,y_test_r):0.2f}')
```



- But now we might be overfitting on the test set!
- How to choose hyperparameters: **Cross-Validation**

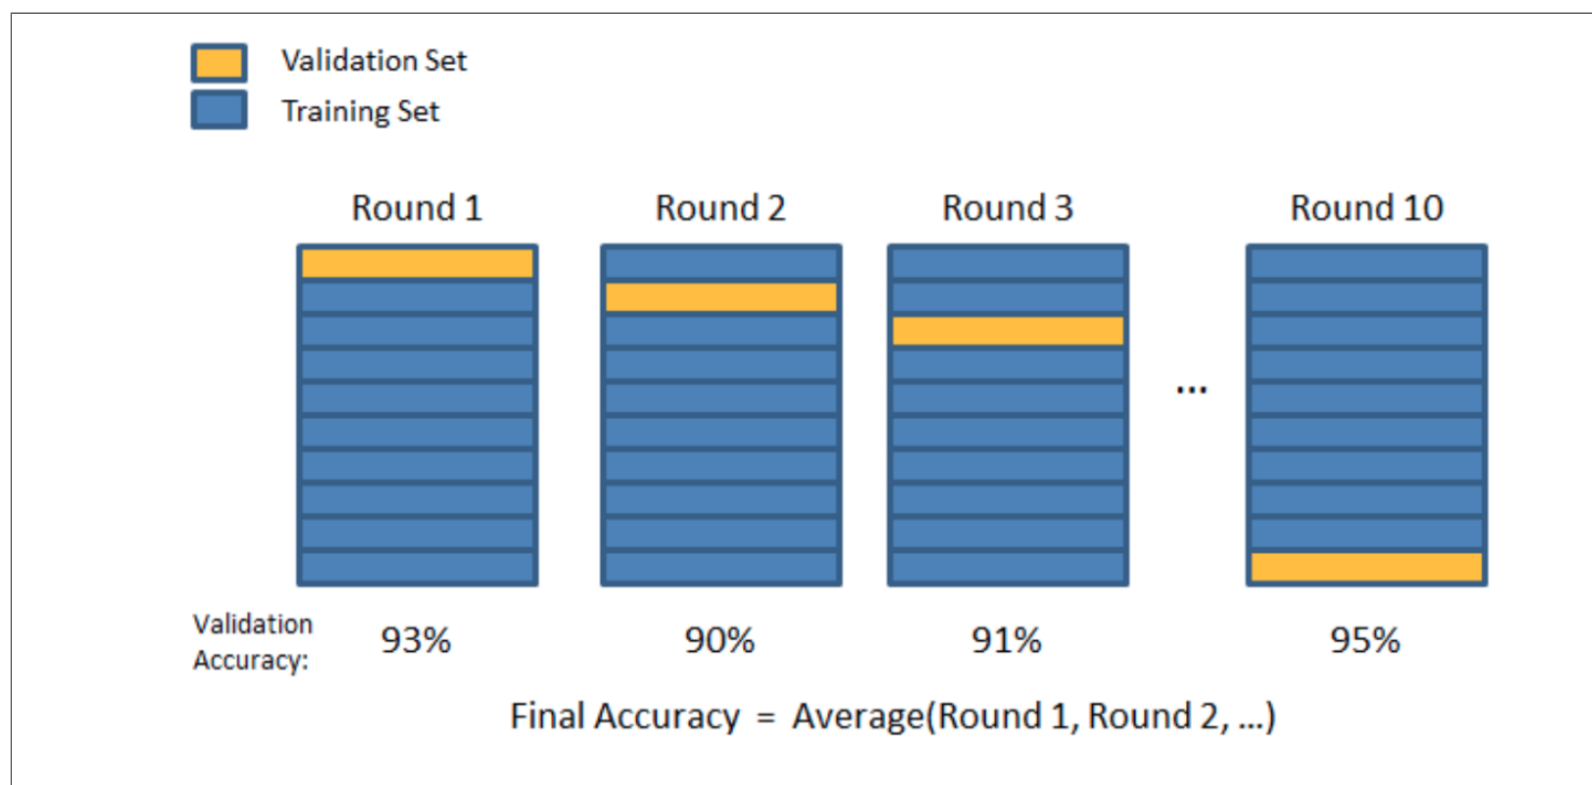
Aside: Hyperparameters

- **parameter:** something learned by the model itself (eg. coefficient in linear model)
- **hyperparameter:** something we set by hand (eg. decision tree max depth)

k -Fold Cross-Validation

1. split dataset into k equal sized subsets (folds)
 2. for each subset (fold)
 - train on the other $k - 1$ subsets combined
 - test on this subset to get a score
 3. average across all scores
- Result is a set of samples of model performance
 - Can use to set hyperparameters without overfitting on train or test
 - Can also use to estimate range of generalization performance

Example: 10-Fold Cross-Validation



k-Fold Cross-Validation Continued

- Can be used for:
 - tuning hyperparameters
 - model selection
 - any time we need estimate of model performance
- **Issue:** each fold requires training the model
 - Training time can be an issue for large k
- What values can k take?
 - min: 2
 - max: n , the size of the dataset (aka Leave-One-Out CV)

k-Fold Cross-Validation in sklearn

In [15]:

```
from sklearn.model_selection import cross_val_score

scores = cross_val_score(DecisionTreeRegressor(max_depth=2),
                          X_train_r,
                          y_train_r,
                          cv=5) #default

scores
```

Out[15]:

```
array([0.41853032, 0.37075208, 0.45405394, 0.2588269 ,
       0.22799476])
```

In [16]:

```
print(f'{np.mean(scores) :0.2f} +- {2*np.std(scores) :0.2f}')
```

```
0.35 +- 0.18
```

Tuning Hyperparameters with CV

In [17]:

```
mean_scores = []

for depth in [1, 2, 3, 5, 10]:
    dtr = DecisionTreeRegressor(max_depth=depth)
    scores = cross_val_score(dtr, X_train_r, y_train_r, cv=5)
    mean_scores.append( (depth, scores.mean().round(3)) )

for depth, mean_score in mean_scores:
    print(f'{depth = :2d} : {mean_score: 0.3f}')
```

```
depth = 1 : 0.289
depth = 2 : 0.346
depth = 3 : 0.341
depth = 5 : 0.050
depth = 10 : -0.142
```

In [18]:

```
# find the depth that gives best score (highest R^2)
sorted(mean_scores, key=lambda x:x[1], reverse=True)[0] # sorted is ascending by default
```


Out[18]:

```
(2, 0.346)
```

Visualize Tuning: Validation Curve

Validation Curve

- Show model complexity vs model performance on both train and test/validation
- Want to find point where performance on validation set begins to decline (overfitting)



From PDSH

Validation Curve in sklearn

In [19]:


```
from sklearn.model_selection import validation_curve

depth = [1,2,3,5,8,10]
train_scores, test_scores = validation_curve(DecisionTreeRegressor(),
                                             X_train_r, y_train_r,
                                             param_name='max_depth',
                                             param_range=depth,
                                             cv=5)

mean_train_scores = np.mean(train_scores,axis=1) # take the mean across columns
mean_test_scores = np.mean(test_scores,axis=1)
```

In [20]:

```
fig,ax = plt.subplots(1,1)
ax.plot(depth, mean_train_scores, 'o-', color='b',label='training score');
ax.plot(depth, mean_test_scores, 'o-', color='r', label='validation score');
ax.set_xlabel('max_depth'), ax.set_ylabel('R^2'); ax.set_title('Validation Curve for DecisionTree');
ax.legend();
```



More Than One HyperParameter? Grid Search

Grid Search: Search over a 'grid' of hyperparameter settings

Example: KNN "number of neighbors" and "distance metric"

In [21]:

```
distance_metrics = ['euclidean', 'manhattan']  
n_neighbors = [1, 3, 5]  
  
grid = []  
for d in distance_metrics:  
    for k in n_neighbors:  
        print([d, k])
```

```
['euclidean', 1]  
['euclidean', 3]  
['euclidean', 5]  
['manhattan', 1]  
['manhattan', 3]  
['manhattan', 5]
```


Grid Search in sklearn

In [22]:

```
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsRegressor

params = {'n_neighbors':[1,2,3,5,10],
          'metric':['euclidean','manhattan']}
gscv = GridSearchCV(KNeighborsRegressor(),
                    param_grid=params,      # grid of size 10
                    cv=3,                  # do 3-fold CV at every grid point
                    refit=True) # refit True trains one more time on the entire training set

gscv.fit(X_train_r,y_train_r)      # How many times are we training a model here? (2*5*3 + 1 = 31)

print(gscv.best_params_)
```

```
{'metric': 'euclidean', 'n_neighbors': 5}
```

In [23]:

```
scores = cross_val_score(gscv.best_estimator_,X_train_r,y_train_r,cv=5)

print(f'{np.mean(scores):0.2f} +- {2*np.std(scores):0.2f}')
```

```
0.35 +- 0.21
```

Review So Far

- Regression Metrics
 - MSE and RMSE
 - R^2
- Model Selection
 - Comparison to Baseline Model
 - Underfitting/Overfitting, Bias/Variance
 - Train/Test Split
- Hyperparameter Tuning
 - Cross-Validation
 - Grid Search
 - Validation Curve

Data Setup for Classification

In [24]:

```
idx_binary = df_wine['class'].isin([0,1])           # reduce to binary classification

X_bc = df_wine.loc[idx_binary,['alcalinity_of_ash','magnesium']]      # only 2 features for ease of plotting
y_bc = df_wine.loc[idx_binary,'class']                                # pull out classification target [0,1]

X_train_bc,X_test_bc,y_train_bc,y_test_bc = train_test_split(X_bc,
                                                                y_bc,
                                                                stratify=y_bc, # maintain label proportions
                                                                random_state=0
                                                                )

pd.DataFrame({'train':y_train_bc.value_counts(),'test':y_test_bc.value_counts()}).sort_index()
```

Out[24]:

	train	test
0	44	15
1	53	18

In [25]:

```
X_mc = df_wine.loc[:,['alcalinity_of_ash', 'magnesium']]      # multiple features for multiclass classification task
y_mc = df_wine.loc[:, 'class']                                # pull out classification target [0,1,2]
X_train_mc,X_test_mc,y_train_mc,y_test_mc = train_test_split(X_mc,
                                                                y_mc,
                                                                stratify=y_mc, # maintain label proportions
                                                                random_state=123
                                                                )

pd.DataFrame({'train':y_train_mc.value_counts(),'test':y_test_mc.value_counts()}).sort_values(by="train")
```

Out[25]:

	train	test
--	-------	------

	train	test
2	36	12
0	44	15
1	53	18

Default Metric in Classification: Accuracy

- **Accuracy:** out of all the observations, how many did I get right?

In [26]:

```
from sklearn.dummy import DummyClassifier
from sklearn.tree import DecisionTreeClassifier
dummyc = DummyClassifier(strategy='prior').fit(X_train_bc,y_train_bc) # works like 'most-frequent'
dtc = DecisionTreeClassifier(max_depth=2).fit(X_train_bc,y_train_bc)

print(f'{dummyc.score(X_test_bc,y_test_bc) = :0.2f}') # default classification score is accuracy
print(f'{dtc.score(X_test_bc,y_test_bc) = :0.2f}')
```

dummyc.score(X_test_bc,y_test_bc) = 0.55

dtc.score(X_test_bc,y_test_bc) = 0.79

- But what if the cost of calling a negative a positive is different from calling a positive a negative?
- Examples:
 - disease testing
 - medical product failures
 - incarceration

Errors in Classification

- There are different kinds of error in classification


		Predicted class	
		P	N
Actual class	P	True positives (TP)	False negatives (FN)
	N	False positives (FP)	True negatives (TN)

From PML

Visualizing Errors with a Confusion Matrix

In [27]:

```
fig,ax = plt.subplots(1,1,figsize=(6,4));
plot_decision_regions(X_train_bc.values,y_train_bc.values,dtc,ax=ax);
sns.scatterplot(x=X_bc.columns[0],y=X_bc.columns[1],data=X_test_bc[y_test_bc == 0],color="blue",s=120);
sns.scatterplot(x=X_bc.columns[0],y=X_bc.columns[1],data=X_test_bc[y_test_bc == 1],color="red",s=120);
```



In [28]:

```
from sklearn.metrics import confusion_matrix

print('training set error\n', confusion_matrix(y_train_bc,dtc.predict(X_train_bc)))
print()
print('test set error\n', confusion_matrix(y_test_bc,dtc.predict(X_test_bc)))
```

training set error
[[41 3]

```
[14 39]]
```

```
test set error
```


```
[[15  0]
```

```
[ 7 11]]
```

Plot Confusion Matrix with mlxtend

In [29]:

```
from mlxtend.plotting import plot_confusion_matrix  
  
fig, ax = plt.subplots(1,1,figsize=(4,4))  
plot_confusion_matrix(confusion_matrix(y_test_bc,dtc.predict(X_test_bc)),axis=ax);
```



Weighing Errors: Precision vs. Recall

Precision

- Out of the observations I predicted positive (TP+FP), how many are truly positive (TP)?

$$\text{precision} = \frac{TP}{TP+FP}$$

Recall

- Out of the truly positive (TP+FN), how many observations did I predict positive (TP)?

$$\text{recall} = \frac{TP}{TP+FN}$$

Using Other Measures in sklearn

In [30]:

```
dummyc_precision_scores = cross_val_score(dummyc,X_train_bc,y_train_bc,cv=5,scoring='precision')
dummyc_recall_scores     = cross_val_score(dummyc,X_train_bc,y_train_bc,cv=5,scoring='recall')

print(f'dummy precision: {np.mean(dummyc_precision_scores):0.2f} +- {2*np.std(dummyc_precision_scores):0.2f}')
print(f'dummy recall    : {np.mean(dummyc_recall_scores):0.2f} +- {2*np.std(dummyc_recall_scores):0.2f}')
print()

dtc_precision_scores = cross_val_score(dtc,X_train_bc,y_train_bc,cv=5,scoring='precision')
dtc_recall_scores     = cross_val_score(dtc,X_train_bc,y_train_bc,cv=5,scoring='recall')

print(f'dtc precision   : {np.mean(dtc_precision_scores):0.2f} +- {2*np.std(dtc_precision_scores):0.2f}')
print(f'dtc recall      : {np.mean(dtc_recall_scores):0.2f} +- {2*np.std(dtc_recall_scores):0.2f}')
```

dummy precision: 0.55 +- 0.04

dummy recall : 1.00 +- 0.00

dtc precision : 0.85 +- 0.36

dtc recall : 0.66 +- 0.38

How do we decide if something is positive or negative?

Usually set a threshold :

$$\hat{y}_i = \begin{cases} 1 & \text{if } P(y_i = 1|x_i) > \text{threshold,} \\ 0 & \text{o.w.} \end{cases}$$

Usually, threshold = .5, but it doesn't have to be.

What happens if we change it?

- High threshold → High Precision, Low Recall
- Low threshold → High Recall, Low Precision

Combining Precision and Recall: F_1 -score

Usually, we just want one number to optimize

F_1 -**score**: harmonic mean of precision and recall

- eg. weighted average of the precision and recall

$$F_1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

Note that F_1 ignores True Negatives!

In [31]:

```
dummyc_f1_scores = cross_val_score(dummyc,X_train_bc,y_train_bc,cv=5,scoring='f1')
dtc_f1_scores = cross_val_score(dtc,X_train_bc,y_train_bc,cv=5,scoring='f1')
print(f'dummyc f1 = {np.mean(dummyc_f1_scores):0.2f} +- {2*np.std(dummyc_f1_scores):0.2f}')
print(f'dtc f1 = {np.mean(dtc_f1_scores):0.2f} +- {2*np.std(dtc_f1_scores):0.2f}')
```

dummyc f1 = 0.71 +- 0.03

dtc f1 = 0.72 +- 0.31

Paying attention to True Negatives: ROC

Receiver Operating Characteristic

- displays FPR vs TPR

$$\text{False Positive Rate (FPR)} = \frac{\text{FP}}{\text{FP} + \text{TN}} = \frac{\text{negatives we got wrong}}{\text{all negatives}}$$

$$\text{True Positive Rate (TPR)} = \text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}} = \frac{\text{positives we got right}}{\text{all positives}}$$

How do these change as we move our threshold?

In [32]:

```
from sklearn.metrics import roc_curve
fpr_dtc, tpr_dtc, _ = roc_curve(y_train_bc, dtc.predict_proba(X_train_bc)[:, -1])
fpr_dtc, tpr_dtc
```

Out[32]:

```
(array([0.          , 0.          , 0.06818182, 1.
]),
 array([0.          , 0.56603774, 0.73584906, 1.
]))
```


Plotting ROC Curves


In [33]:

```
def plot_roc(curves):
    fig, ax = plt.subplots(1, 1, figsize=(6, 6))
    lw = 2
    for fpr, tpr, model_name in curves:
        l1, = ax.plot(fpr, tpr, lw=lw, label=model_name)
    ax.plot([0, 1], [0, 1], color='k', lw=lw, linestyle='--')
    ax.set_xlim([0.0, 1.0])
    ax.set_ylim([0.0, 1.05])
    ax.set_xlabel('False Positive Rate')
    ax.set_ylabel('True Positive Rate')
    ax.set_aspect('equal', 'box')
    ax.set_title('Receiver operating characteristic example')
    ax.legend()
```

Plotting ROC Curves

In [34]:


```
curves = [(fpr_dtc, tpr_dtc, 'dtc2')]
fpr_dummy, tpr_dummy, _ = roc_curve(y_train_bc, dummyc.predict_proba(X_train_bc)[:,-1]) # Compare dummy
curves.append((fpr_dummy, tpr_dummy, 'dummy'));
plot_roc(curves);
```



Plotting ROC Curves

In [35]:

```
for depth in [3,4]:  
    fpr, tpr, _ = roc_curve(y_train_bc, DecisionTreeClassifier(max_depth=depth).fit(X_train_bc, y_train_bc).predict_proba(X_train_bc)[:,1])  
    curves.append((fpr, tpr, 'dtt'+str(depth)))  
plot_roc(curves);
```



ROC AUC

- But again, we'd like one number to optimize
- ROC **A**rea **U**nder the **C**urve
 - How much area falls under the ROC curve?

In [36]:

```
dummyc_rocauc_scores = cross_val_score(dummyc,X_train_bc,y_train_bc,cv=5,scoring='roc_auc')
dtc_rocauc_scores = cross_val_score(dtc,X_train_bc,y_train_bc,cv=5,scoring='roc_auc')

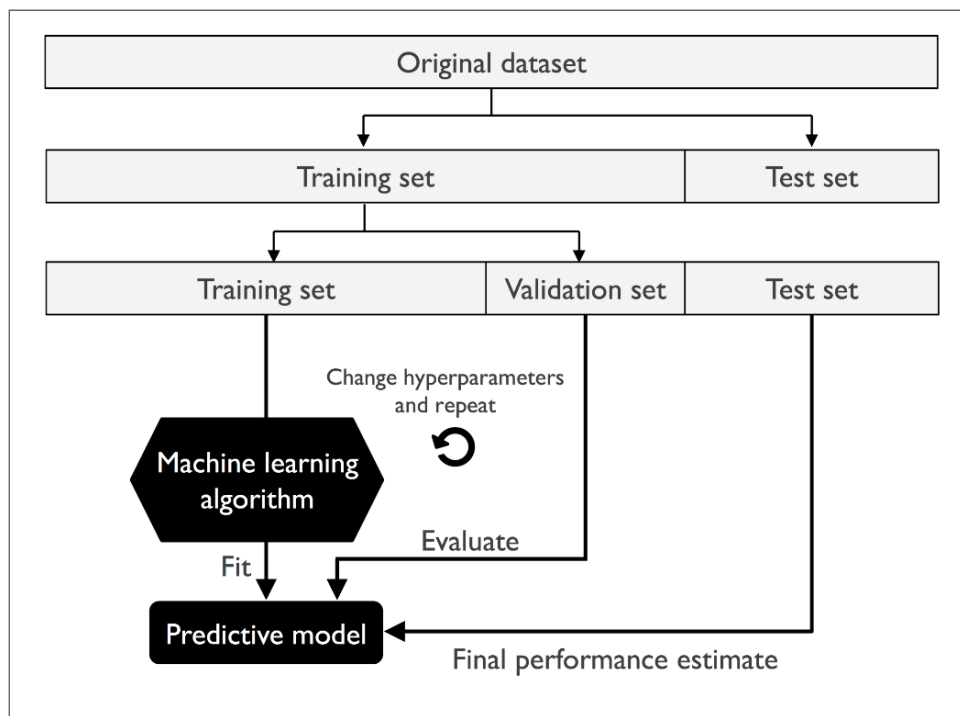
print(f'dummyc rocauc = {np.mean(dummyc_rocauc_scores):0.2f} +- {2*np.std(dummyc_rocauc_scores):0.2f}')
print(f'dtc rocauc = {np.mean(dtc_rocauc_scores):0.2f} +- {2*np.std(dtc_rocauc_scores):0.2f}')
```

dummyc rocauc = 0.50 +- 0.00

dtc rocauc = 0.78 +- 0.22

Review: Steps to Choosing a Model


1. Create Held-Aside Set (Train/Test Split)
2. Determine Metric (or combination of metrics)
3. Get a Baseline for comparison
4. Use Cross-Validation to fit Hyperparameters and Choose Model
5. Evaluate Chosen Model on Held-Aside Set



From PML

Avoiding Overfitting in Linear Models: Regularization


- Use to avoid overfitting in linear models
- Idea: can we reduce complexity of our linear model by minimizing weights?



From **[https://www.wikiwand.com/en/Regularization \(mathematics\)](https://www.wikiwand.com/en/Regularization_(mathematics))**

Regression: Finding the Weights

- Linear models learn by finding weights that minimize a cost.
- Can we get close to the solution while still keeping weights small (simpler model)?




From PML

Penalizing extreme weights (w)

$$\arg \min_w C(f(w, x), y)$$
$$\arg \min_w C(f(w, x), y) + \lambda g(w)$$

Regularization: Ridge


- Coefficients are encouraged to be "small"
- Uses L_2 norm: $\|w\|_2 = \sqrt{\sum_{j=1}^m w_j^2}$



From PML

Regularization: LASSO

- Coefficients are (likely) driven to zero
- Uses L_1 norm: $\|w\|_1 = \sum_{j=1}^m |w_j|$



From PML

Regularization: ElasticNet

- Mixture of L_1 and L_2
- $\alpha L_1 + (1 - \alpha)L_2$
- introduces a new hyperparameter α or `l1_ratio`
- `l1_ratio` = 1 is LASSO (L_1)
- `l1_ratio` = 0 is Ridge (L_2)

Regularization in sklearn

In [37]:

```
from sklearn.linear_model import LogisticRegression

logr = LogisticRegression(penalty='l2', # default
                          C=1.0,      # weight on regularization, 1/lambda above
                          l1_ratio=None # only used when penalty is 'elasticnet'
                          )
```

In [38]:

```
for C in [.001, .1, 1, 10, 1000]:
    logr = LogisticRegression(penalty='l2', # default
                              C=C,         # weight on regularization, 1/lambda above
                              ).fit(X_train_bc, y_train_bc)
    print(f'{str(C):5s} : {logr.coef_[0].round(2)}')
```

```
0.001 : [ 0.02 -0.02]
0.1    : [ 0.69 -0.57]
1      : [ 1.09 -0.93]
10     : [ 1.19 -1.02]
1000   : [ 1.2  -1.03]
```

In [39]:

```
for C in [.001, .1, 1, 10, 1000]:
    logr = LogisticRegression(penalty='l1',
                              C=C,         # weight on regularization, 1/lambda above
                              solver='liblinear'
                              ).fit(X_train_bc, y_train_bc)
    print(f'{str(C):5s} : {logr.coef_[0].round(2)}')
```

0.001	:	[0. 0.]
0.1	:	[0.5 -0.37]
1	:	[1.08 -0.92]
10	:	[1.19 -1.02]
1000	:	[1.2 -1.03]

GridSearchCV with Regularization

In [40]:

```
param_grid = {'l1_ratio':[0,.5,1],
              'C': [.001,.01,1,10]}
logr_gscv = GridSearchCV(estimator=LogisticRegression(penalty='elasticnet', solver='saga'),
                        param_grid=param_grid,
                        cv=3,
                        n_jobs=-1).fit(X_train_bc, y_train_bc)

print(f'best parameter setting found: {logr_gscv.best_params_}')
print(f'best coefficients found      : {logr_gscv.best_estimator_.coef_[0]}')
print(f'best training score found   : {logr_gscv.best_score_.round(3)}')

logr_gscv_test_score = logr_gscv.score(X_test_bc, y_test_bc)
logr_noreg_test_score = (LogisticRegression(penalty='none')
                        .fit(X_train_bc, y_train_bc)
                        .score(X_test_bc, y_test_bc)
                        )

print()
print(f'logr_gscv test score  : {logr_gscv_test_score.round(3)}')
print(f'logr_noreg test score : {logr_noreg_test_score.round(3)}')
```

```
best parameter setting found: {'C': 1, 'l1_ratio': 1}
best coefficients found      : [ 1.09869298 -0.9278305
 8]
best training score found   : 0.825
```

```
logr_gscv test score  : 0.818
logr_noreg test score : 0.818
```


ElasticNetCV

In [41]:

```
from sklearn.datasets import make_regression
from sklearn.linear_model import ElasticNetCV

X_synth,y_synth = make_regression(n_samples=100,
                                n_features=200,
                                n_informative=10,
                                random_state=123
                                )

X_synth_train,X_synth_test,y_synth_train,y_synth_test = train_test_split(X_synth, y_synth, random_state=123)

dummy_synth = DummyRegressor(strategy='mean').fit(X_synth_train,y_synth_train)
lr_synth = LinearRegression().fit(X_synth_train,y_synth_train)
en_synth = ElasticNetCV(alphas=[.01,.1,1,100]).fit(X_synth_train,y_synth_train)

print(f'found alpha: {en_synth.alpha_}, found l1_ratio: {en_synth.l1_ratio_}\n')
print(f'{dummy_synth.score(X_synth_train,y_synth_train)} = : 0.2f}')
print(f'{lr_synth.score(X_synth_train,y_synth_train)} = : 0.2f}')
print(f'{en_synth.score(X_synth_train,y_synth_train)} = : 0.2f}\n')
print(f'{dummy_synth.score(X_synth_test,y_synth_test)} = : 0.2f}')
print(f'{lr_synth.score(X_synth_test,y_synth_test)} = : 0.2f}')
print(f'{en_synth.score(X_synth_test,y_synth_test)} = : 0.2f}')
```

found alpha: 1.0, found l1_ratio: 0.5

```
dummy_synth.score(X_synth_train,y_synth_train) = 0.00
lr_synth.score(X_synth_train,y_synth_train) = 1.00
en_synth.score(X_synth_train,y_synth_train) = 0.95
```

```
dummy_synth.score(X_synth_test,y_synth_test) = -0.00  
lr_synth.score(X_synth_test,y_synth_test) = 0.16  
en_synth.score(X_synth_test,y_synth_test) = 0.24
```

Questions?